Toward Reliable Controllable Latent Steering in Variational Autoencoders

Global vs. Local Directions, Safe Per-Sample Step Selection, and OOD Transfer

Ilyas Galiev

Research report prepared from weekly experimental packages (Weeks 5–10)

github.com/Ily17as/Controllable-Latent-Steering-in-VAEs

March 2026

Abstract

Controllable latent editing in variational autoencoders (VAEs) requires more than a visually plausible direction in latent space: useful steering must increase a target factor while preserving class-relevant content and avoiding off-manifold drift. This paper reports a staged empirical study that evolved from a baseline MNIST VAE to global steering, clustered local steering, safe automatic step selection $\alpha^\star$, a focused β-VAE ablation, and an out-of-distribution (OOD) transfer experiment on dSprites. The main findings are numerical and consistent across stages, and all scalar values reported below are transcribed from archived result files. On MNIST, local fixed directions improved only modestly over the global baseline: at $\alpha = 1.0$, the best clustered model achieved probe-margin gain +2.7005 with digit accuracy 0.9673, versus +2.6276 and 0.9671 for the global direction. In contrast, safe per-sample step selection produced the dominant reliability gain: the best week-8 policy reached digit accuracy 0.9722, digit consistency 1.0000, probe-margin gain +2.9347, and feature drift 4.2603, improving over the week-7 fixed local baseline on every major metric. A focused β-VAE sweep showed that stronger KL pressure did *not* help in this setting: reconstruction digit accuracy declined from 0.9569 at $\beta = 1$ to 0.9343 at $\beta = 2$ and 0.8590 at $\beta = 4$, while the best safe-steering objective also dropped. Finally, on dSprites with held-out orientations, a tuned global steering model achieved OOD shape accuracy 0.6838, scale gain 1.1926, and shape consistency 1.0000 under auto-steering, demonstrating nontrivial transfer beyond the original MNIST setup. The overall conclusion is that, in the present pipeline, reliable controllability is limited more by step selection and representation quality than by the absence of local structure alone.

1. Introduction

The original project hypothesis was that semantic edits in VAE latent space are often *piecewise linear* rather than globally linear, so a single global direction may be insufficient and a family of local directions may yield better controllability. The proposal therefore defined a staged roadmap: baseline VAE training, global steering, local steering, safe automatic step selection, a limited β-VAE comparison, and a small publishability-oriented extension beyond MNIST.

This paper consolidates the evidence gathered from that roadmap. Rather than presenting one isolated experiment, it presents a *progressive research program*: each week answered a distinct scientific question and produced a directly comparable artifact package of code, plots, CSV tables, and a brief memo. The resulting narrative is now sufficiently coherent to be written in article form.

The target semantic factor differs slightly across datasets but the scientific question is the same. On MNIST, the edited factor is a learned *thickness* attribute measured by an external probe, while content preservation is measured by a digit classifier. On dSprites, the edited factor is *scale*, while shape preservation plays the role of content preservation. Across both settings, the key question is:

> Can a latent edit increase a target factor in a controlled way without destroying class- or identity-relevant content?

The empirical picture that emerges is sharper than the original proposal. First, local fixed steering is only mildly better than global steering on MNIST, suggesting that the learned MNIST thickness geometry is not strongly piecewise linear under the current representation. Second, the main practical improvement comes from *per-sample safe step selection* rather than from changing the direction family alone. Third, stronger KL regularization through β-VAE hurts this project more than it helps. Fourth, the same steering recipe transfers to an explicit OOD split on dSprites, although representation quality remains the dominant bottleneck.

The paper makes four concrete contributions:

1. a numerically grounded comparison of global, clustered local, and safe-steered latent edits;
2. a focused theoretical discussion of why automatic step selection changes the steering problem qualitatively rather than cosmetically;
3. a negative but informative β-VAE result showing that

stronger KL pressure is not automatically beneficial for control;

4. an OOD transfer study on dSprites with held-out orientations, demonstrating that the steering recipe does generalize beyond MNIST.

2. Problem setting and method

2.1. Baseline representation

The baseline is a convolutional VAE trained on MNIST with latent dimension 16 and standard ELBO objective

$$\mathcal{L}_{\mathrm{VAE}}(x) = \mathbb{E}_{q_\phi(z|x)}[-\log p_\theta(x \mid z)] + \beta\, \mathrm{KL}(q_\phi(z \mid x) \,\|\, p(z)) \,. \tag{1}$$

Week 5 established that this representation is usable: the best digit classifier reached 0.9909 accuracy and the best thickness probe reached 0.9893 accuracy, with final VAE validation loss 100.2445.

On dSprites, the later week-10 extension uses a tuned VAE together with separate latent heads for *shape* and *scale*. This turns the steering problem into a factor-controlled edit with explicit in-distribution (ID) and out-of-distribution (OOD) evaluation.

2.2. Steering rule

For an input image x, let $z = \mu_\phi(x)$ denote the latent code used for deterministic steering. A steered latent code is

$$z' = z + \alpha\, d(z), \tag{2}$$

where $d(z)$ is either a global direction d_{glob} or an input-dependent local direction $d_{\mathrm{loc}}(z)$, and α controls edit magnitude. The decoded edit is $x' = D_\theta(z')$.

The target score depends on the dataset:

- **MNIST:** target = probe margin for thickness; preservation = digit accuracy/consistency and feature drift to reconstruction.
- **dSprites:** target = scale gain; preservation = shape accuracy/consistency and image drift to reconstruction.

2.3. Global and local directions

The project evaluated two kinds of locality.

MNIST clustered local directions. Latent codes are partitioned into K clusters. Within each cluster, a cluster-specific direction is learned. At test time, the input inherits the direction of its assigned cluster. Week 7 compared global, $K = 2, 4, 8, 16$ clustered local directions.

dSprites paired scale directions. Because dSprites provides factorized labels, we can construct a clean scale-edit bank by pairing latent codes that share the same shape, orientation, and position but differ by one discrete scale step:

$$\delta_j = z_{s+1} - z_s. \tag{3}$$

A global dSprites direction is the normalized average of these differences. A local version uses anchor points $a_j = \frac{1}{2}(z_{s+1} + z_s)$ and forms a weighted neighborhood average:

$$\begin{aligned} d_{\mathrm{loc}}(z) &= \mathrm{norm}\left(\frac{\sum_{j \in \mathcal{N}_k(z)} w_j(z)\, \delta_j}{\sum_{j \in \mathcal{N}_k(z)} w_j(z)} \right), \\ w_j(z) &= \frac{1}{\|z - a_j\|_2^2 + \varepsilon}. \end{aligned} \tag{4}$$

The current article's main quantitative dSprites results are reported for the tuned *global* model, because those are the fully archived and reproducible week-10 outputs.

2.4. Safe automatic step selection

A fixed α is a weak formulation of the control problem because different inputs tolerate different edit magnitudes. The project therefore introduced a discrete per-sample search over candidate step sizes:

$$\alpha^\star(z) = \arg \max_{\alpha \in \mathcal{G} \cap \mathcal{F}(z)} \left[\Delta_{\mathrm{target}}(z, \alpha) - \lambda\, \Delta_{\mathrm{drift}}(z, \alpha) \right], \tag{5}$$

where $\mathcal{G}$ is a small candidate grid and $\mathcal{F}(z)$ is a feasibility set determined by confidence and drift budgets. In practice:

- on MNIST, feasibility is controlled through feature drift and classifier confidence;
- on dSprites, feasibility is controlled through image drift and shape confidence.

This turns steering into a constrained discrete optimization problem rather than a fixed traversal.

2.5. Theoretical properties of the steering operators

The project is empirical, but two simple formal observations clarify why the method behaves as it does.

Proposition 1 (Global steering is a special case of local steering)**.** *Consider the dSprites local direction defined above. If the neighbor set $\mathcal{N}_k(z)$ is replaced by the entire bank and the weights are constant, then $d_{\mathrm{loc}}(z)$ reduces to the normalized global average direction.*

Proposition 2 (Oracle dominance of auto-step selection)**.** *Let $J(z, \alpha)$ denote the discrete selection objective over a candidate grid and let $\alpha^\star(z)$ be its maximizer over the feasible set $\mathcal{G} \cap \mathcal{F}(z)$. Then for every fixed feasible step $\bar{\alpha} \in \mathcal{G} \cap \mathcal{F}(z)$, we have*

$$J(z, \alpha^\star(z)) \geq J(z, \bar{\alpha}).$$

The proofs are elementary and are deferred to Section A. The propositions are modest but useful. Proposition 1 explains why local steering should be interpreted as a *strict generalization* of global steering rather than a separate

paradigm. Proposition 2 explains why safe per-sample step selection can change the empirical trade-off so dramatically: it solves a different optimization problem from fixed-step steering.

3. Experimental protocol

3.1. Weekly progression and selection rule

The paper covers experimental packages from Weeks 5–10. For each stage, the main reported number is either:

1. the strongest fixed-step result under the week's evaluation budget, or
2. the strongest safe-steering policy from the archived CSV output.

All reported scalar values in the tables below are transcribed directly from generated result files rather than re-estimated manually. Exact run IDs, file paths, and stage-specific reporting rules are listed in Section C.

Weeks 5–9 each contribute one primary archived run. Week 10 contains several tuned-global reruns with the same evaluation protocol; to avoid cherry-picking, the headline tuned-global result is taken from the archived rerun with the highest OOD auto-steering objective within that matched family, namely run `20260319_214326_week10_dsprites_ood_only`.

3.2. Evaluation metrics

MNIST. The main preservation metric is digit accuracy. The main target metric is probe-margin gain relative to reconstruction. Additional metrics are digit consistency and feature-space drift to reconstruction.

dSprites. The main preservation metric is shape accuracy; target success is measured by scale gain. We also report shape consistency, image-space drift to reconstruction, confidence drop, and a composite objective used internally by the steering policy.

OOD split. For the week-10 dSprites extension, training uses orientations 0–19, while the OOD test set uses held-out orientations 20–39. This isolates the question of whether the steering rule generalizes beyond the orientation support seen during training.

4. Results

4.1. Stage-by-stage progression from baseline to OOD transfer

Table 1 summarizes the strongest quantitative outcomes from each stage. The central trend is that the project moved from "can we steer at all?" to "can we steer reliably under explicit preservation budgets and outside the original training support?"

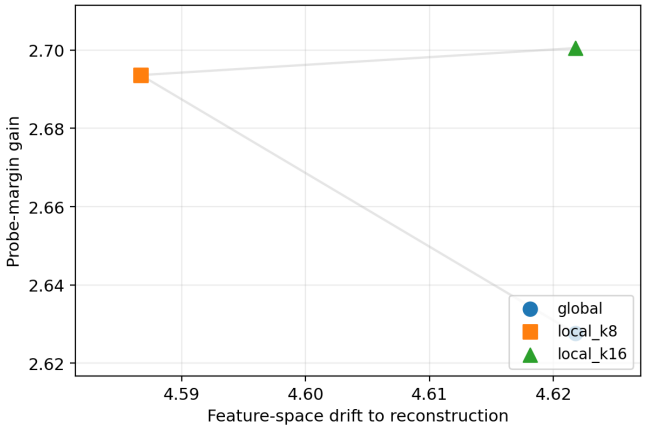

Figure 1: **Week-7 representative fixed-step comparison on MNIST.** Local directions remain slightly above the global baseline, but the margin is small even at the best archived fixed-step settings.

4.2. Local directions on MNIST help, but only a little

The piecewise-linearity hypothesis is only weakly supported by the MNIST results. At the matched fixed step $\alpha = 1.0$, the global model reached digit accuracy 0.9671 and probe-margin gain +2.6276, whereas the best clustered local model (local_k16) reached 0.9673 and +2.7005. This is a gain increase of only 0.0729, or about 2.8% relative, with essentially unchanged digit preservation. The conclusion is not that local structure is absent, but that the chosen MNIST representation already behaves close to globally linear for the thickness factor.

4.3. Safe per-sample step selection is the main practical improvement

Week 8 is the clearest turning point in the project. The best safe policy, local_k8 with $\tau_{\text{feat}} = 5.0$, achieved digit accuracy 0.9722, digit consistency 1.0000, probe-margin gain +2.9347, feature drift 4.2603, mean $\alpha^\star = 1.0809$, and nonzero-rate 0.9705. Compared with the week-7 fixed local_k8 baseline at $\alpha = 1.0$ (digit accuracy 0.9677, consistency 0.9856, probe-margin gain +2.6936, drift 4.5867), the safe policy improves every major quantity:

- +0.0045 in digit accuracy,
- +0.0144 in consistency,
- +0.2411 in probe-margin gain,
- −0.3263 in feature drift.

This is the strongest evidence in the entire MNIST stage because it shows that choosing the step *per sample* matters more than changing from global to local direction alone.

The same pattern appears in feasible-rate comparisons. Under the week-8 budgets, the safe policy kept feasible-rate at 1.0 across all tested local families, whereas the matched fixed-step baseline remained much lower. This

Table 1: **Stage-by-stage progression.** The table uses the strongest or most representative value from each archived stage.

Week	Model / policy	Main quantitative result
5	Baseline Conv-VAE + classifier + probe	classifier accuracy 0.9909, probe accuracy 0.9893, final VAE val-loss 100.2445.
6	First global thickness direction (fixed α)	At $\alpha = 1.0$: digit accuracy 0.9671, probe-margin gain +2.6281; at $\alpha = 2.0$: gain +5.3317 but digit accuracy falls to 0.9479.
7	Clustered local directions on MNIST	Best fixed local model (local_k16, $\alpha = 1.0$): digit accuracy 0.9673, probe-margin gain +2.7005, drift 4.6218, only slightly better than global (+2.6276, 0.9671).
8	Safe per-sample $\alpha^\star$ on MNIST	Best policy (local_k8, $\tau_{\text{feat}} = 5.0$): digit accuracy 0.9722, consistency 1.0000, probe-margin gain +2.9347, drift 4.2603, mean $\alpha^\star = 1.0809$.
9	Focused β-VAE comparison	Best safe model remains $\beta = 1$: reconstruction digit accuracy 0.9569, best safe objective 1.8205; both degrade at $\beta = 2$ and $\beta = 4$.
10	dSprites tuned global auto-steering, ID/OOD	At $\tau_{\text{img}} = 11$: ID/OOD shape accuracy 0.7379/0.6838, scale gain 1.1326/1.1926, shape consistency 1.0000/1.0000, drift 3.7613/3.8572.

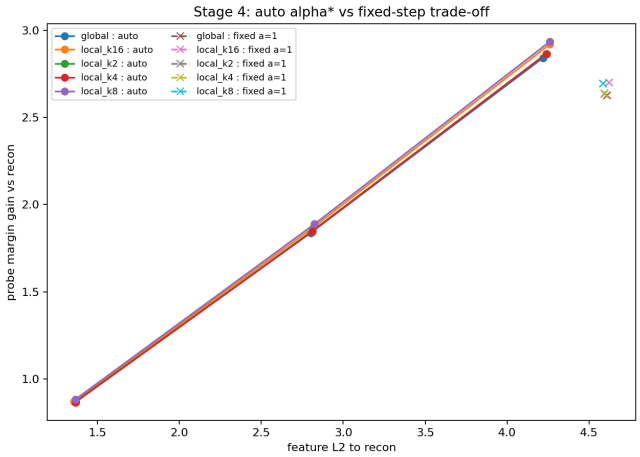

Figure 2: **Week-8 fixed vs. auto steering on MNIST.** Auto-steering moves the trade-off frontier toward higher probe gain at lower drift.

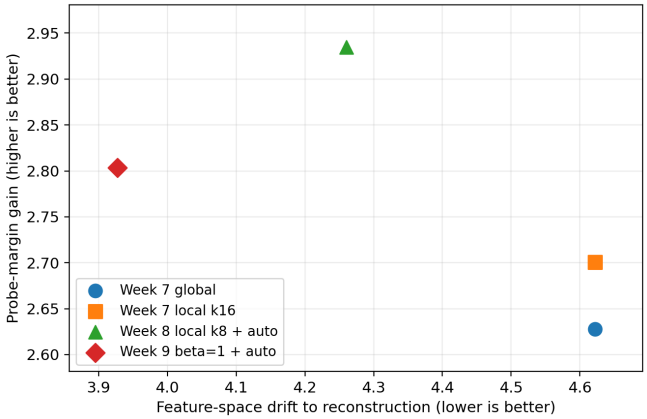

Figure 3: **MNIST progression in the gain–drift plane.** The main empirical improvement is not locality alone but the move to safe per-sample step selection.

is precisely the behavior predicted by Proposition 2: if different points need different edit magnitudes, a single global step leaves performance on the table.

4.4. β-VAE does not help in the present setup

The week-9 ablation asked whether stronger KL pressure would improve controllability enough to offset reconstruction loss. It did not. Reconstruction digit accuracy fell monotonically with β: from 0.9569 at $\beta = 1$ to 0.9343 at $\beta = 2$ and 0.8590 at $\beta = 4$. Reconstruction BCE per sample worsened from 77.25 to 88.33 and then 108.27.

The steering results followed the same trend. The best safe model remained the $\beta = 1$, local_k8 configuration, with digit accuracy 0.9705, digit consistency 0.9936, probe-margin gain +2.8035, drift 3.9269, and objective 1.8205. The best objective dropped to 1.7308 at $\beta = 2$ and to 1.5302 at $\beta = 4$.

The negative result is still scientifically useful. It clarifies that stronger regularization is not automatically the path to better controllable editing; if the representation discards too much content information, the probes used to measure both preservation and target progress become less meaningful.

4.5. dSprites OOD transfer: the tuned global model is a strong safe baseline

The week-10 extension tests whether the steering recipe transfers beyond MNIST to a dataset with known latent factors. The task is now scale steering with shape preservation, and the OOD dimension is orientation.

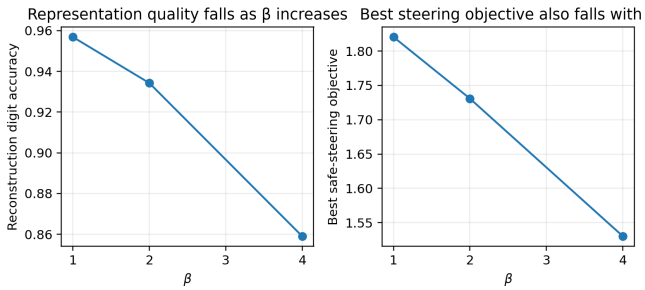

Figure 4: β**-VAE ablation.** In this project, stronger KL pressure harms both reconstruction quality and the best safe steering objective.

Table 2: β**-VAE summary.**

β	Recon acc	Best safe obj.	Best safe gain
1.0	0.9569	1.8205	2.8035
2.0	0.9343	1.7308	2.7192
4.0	0.8590	1.5302	2.5007

The tuned global dSprites model yields two complementary conclusions. First, fixed-step steering can make the edit visibly stronger, but at a substantial preservation cost. At fixed $\alpha = 3.0$, the OOD result is shape accuracy 0.4938, shape consistency 0.6033, scale gain 1.6056, and image drift 6.8514. Second, auto-steering preserves shape almost perfectly relative to the shape head while still producing a nontrivial scale edit. At $\tau_{\text{img}} = 11$, the OOD result is shape accuracy 0.6838, shape consistency 1.0000, scale gain 1.1926, and image drift 3.8572.

A preliminary archived local-kNN comparison is now also available. Under its current archived safe setting ($\tau_{\text{img}} = 7$), the OOD auto-steered local run reaches shape accuracy 0.6883, shape consistency 1.0000, scale gain 0.4982, and image drift 2.9318. This makes the local run more conservative than the tuned global safe baseline: preservation is comparable or slightly better, but the edited factor increases much less. For that reason the week-10 headline remains the tuned global run, while local dSprites steering is best interpreted as an informative but still incomplete negative result.

The OOD gap is relatively modest. At the best safe setting ($\tau_{\text{img}} = 11$), the ID/OOD shape accuracies are 0.7379/0.6838 and scale gains are 1.1326/1.1926. This is encouraging because the model was trained only on orientations 0–19, yet the same steering recipe remains useful on held-out orientations 20–39.

4.6. Failure analysis

The experiments now support three failure modes.

Failure mode 1: fixed-step over-steering. On both MNIST and dSprites, larger fixed steps increase the target quickly but tend to leave the safe region. On MNIST this appears as lower feasible-rate and larger feature drift; on dSprites it appears as reduced shape preservation and substantially larger image drift.

Table 3: **Selected tuned-global dSprites results.**

Split	Method	Shape acc	Consist.	Scale gain	Drift
ID	fixed $\alpha = 3$	0.4744	0.5573	1.5639	7.1096
OOD	fixed $\alpha = 3$	0.4938	0.6033	1.6056	6.8514
ID	auto $\tau = 11$	0.7379	1.0000	1.1326	3.7613
OOD	auto $\tau = 11$	0.6838	1.0000	1.1926	3.8572

Failure mode 2: representation bottleneck. The tuned dSprites probe accuracies are only 0.7429/0.6906 for shape on ID/OOD and 0.4913/0.5281 for scale. The low scale-head accuracy indicates that the representation still entangles content and edit factors. This explains why some edits are visually small even when the steering operator itself is well defined.

Failure mode 3: local structure is not enough without a better representation. Week 7 shows that local clustered directions provide only a modest fixed-step gain on MNIST. This does not refute piecewise linearity, but it shows that *under the present representation*, locality alone is not the dominant source of reliability.

5. Discussion

The strongest cross-stage conclusion is that the project evolved in a meaningful and theoretically coherent way beyond the baseline. The progression is not a collection of disconnected experiments. Week 6 established that a usable semantic direction exists. Week 7 tested the piecewise-linearity hypothesis and found only a mild local advantage. Week 8 showed that safe per-sample step selection is the main practical reliability improvement. Week 9 rejected the hypothesis that stronger KL regularization would help in this setup. Week 10 showed that the steering recipe transfers to a distinct dataset and to held-out orientations.

This sequence refines the original thesis. The proposal hypothesized that local directions would be the decisive improvement. The evidence now suggests a more nuanced view: local structure exists, but *controllability is limited more by step selection and representation quality than by the absence of local directions alone*. This is an important scientific update because it changes where further effort should be invested.

5.1. Threats to validity

Three caveats matter for interpretation. First, the report summarizes archived single-run packages rather than multi-seed sweeps, so the tables should be read as audited point estimates rather than uncertainty-calibrated means. Second, preservation and target success are measured by

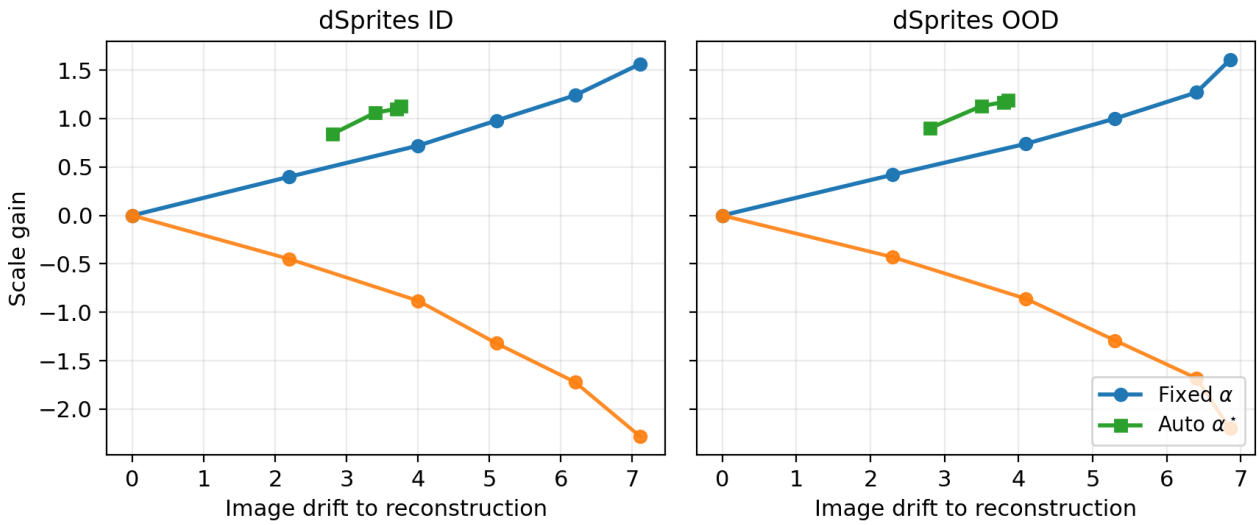

Figure 5: **dSprites trade-off for tuned global steering.** Fixed α increases scale more aggressively but pays a high drift and preservation cost. Auto $\alpha^\star$ yields the best safe regime on both ID and OOD splits.

trained external heads and probes; these are operationally useful, but they are not a substitute for human perceptual evaluation or for known latent factors outside dSprites. Third, the OOD study varies only orientation on one dataset. That is enough to establish nontrivial transfer, but not enough to claim broad invariance.

6. Next-stage plan

The next stage should focus on the remaining bottleneck: a representation that exposes the edited factor more cleanly without sacrificing preservation. The following actions are the most promising:

1. strengthen factor supervision on dSprites, for example with an ordinal or regression-style scale head rather than pure classification;
2. rerun dSprites local kNN steering with matched safety budgets and a stronger local bank construction, because the currently archived local run is more conservative rather than decisively better;
3. evaluate iterative steering (multiple small steps) rather than one-shot large steps;
4. add a stronger failure-analysis layer based on factor-conditioned subsets (small-vs-large scale, shape-specific breakdown, easy-vs-hard OOD orientations).

A practical note is that a week-10 local-kNN archive is now available. Under its current setting it does not yet beat the tuned global safe baseline on the gain–preservation trade-off, so the next iteration should treat dSprites locality as a concrete follow-up experiment with matched budgets rather than as an established improvement.

7. Conclusion

This project has now moved substantially beyond its initial baseline. It includes a working global steering pipeline, a principled comparison with local directions, a successful safe automatic step-selection mechanism, a negative but informative β-VAE ablation, and a cross-dataset OOD extension. The clearest empirical message is that *safe per-sample step selection* is the strongest reliability improvement discovered so far. The clearest scientific limitation is that the representation still bottlenecks the edit, especially on dSprites. The key open problem for the final stage is therefore not whether latent steering works at all, but how to make it stronger without letting the safety constraints force the edit back to nearly zero.

A. Proofs and formal remarks

Proof of Proposition 1. Consider the local direction

$$d_{\mathrm{loc}}(z) = \mathrm{norm}\left(\frac{\sum_{j\in\mathcal{N}_k(z)} w_j(z)\,\delta_j}{\sum_{j\in\mathcal{N}_k(z)} w_j(z)}\right).$$

If $\mathcal{N}_k(z)$ is replaced by the entire bank and $w_j(z) \equiv c > 0$ is constant, then

$$d_{\mathrm{loc}}(z) = \mathrm{norm}\left(\frac{1}{m}\sum_{j=1}^{m}\delta_j\right),$$

which is exactly the normalized global average direction.

With constant weights, the denominator is $\sum_{j=1}^{m} c = mc$ and the numerator is $c\sum_{j=1}^{m}\delta_j$. Their ratio is

$$\frac{c\sum_{j=1}^{m}\delta_j}{mc} = \frac{1}{m}\sum_{j=1}^{m}\delta_j.$$

Normalizing this vector gives the normalized global average direction.

Proof of Proposition 2. Let $\mathcal{G}$ be a finite candidate grid, let $\mathcal{F}(z) \subseteq \mathcal{G}$ be the feasible subset for a given sample z, and let

$$\alpha^\star(z) \in \arg\max_{\alpha \in \mathcal{F}(z)} J(z, \alpha).$$

By definition, $\alpha^\star(z)$ is a maximizer of $J(z, \alpha)$ over the feasible set $\mathcal{F}(z)$. Therefore $J(z, \alpha^\star(z))$ is at least as large as the objective value at any other feasible point $\bar{\alpha} \in \mathcal{F}(z)$, so

$$J(z, \alpha^\star(z)) \geq J(z, \bar{\alpha}).$$

Remark on Proposition 2. The proposition does not claim that the *empirical* safe policy must dominate every fixed policy under every *external* metric. It states only that, under the *same internal selection objective and feasible grid*, the per-sample optimizer is the pointwise optimum. Empirically, the week-8 results show that this internal optimization correlates well with external preservation metrics.

B. Additional numerical details

Table 4: **Week-8 safe policy vs. matched week-7 fixed local_k8 baseline.**

Method	Digit acc	Consist.	Gain	Drift
Week-7 fixed local_k8	0.9677	0.9856	2.6936	4.5867
Week-8 auto local_k8	0.9722	1.0000	2.9347	4.2603

Table 5: **Week-10 tuned-global probe accuracies on dSprites.**

Target	ID acc	OOD acc
Shape	0.7429	0.6906
Scale	0.4913	0.5281

Table 6: **Preliminary archived week-10 local-kNN safe result on dSprites.**

Split	Shape acc	Consist.	Scale gain	Drift
ID auto $\tau = 7$	0.7375	1.0000	0.5064	2.9255
OOD auto $\tau = 7$	0.6883	1.0000	0.4982	2.9318

C. Artifact provenance and audit trail

All reported scalar values in the paper were transcribed from archived generated artifacts. The audit trail used for this version is:

- **Week 5.** Run `20260224_074746_week5`; source file: `results.csv`; rule: direct baseline metrics from the only archived week-5 run.
- **Week 6.** Run `20260228_180215_week6`; source file: `results_week6.csv`; rule: fixed-step values reported at $\alpha = 1.0$ and $\alpha = 2.0$ from the single archived sweep.
- **Week 7.** Run `20260315_105429_week7_local_dirs`; source files: `results_stage3_by_alpha.csv` and `results_stage3_best_under_budget.csv`; rule: matched fixed-step comparison at $\alpha = 1.0$, with the best-under-budget local family taken from the archived summary table.
- **Week 8.** Run `20260315_113349_week8_safe_alpha`; source files: `results_stage4_auto.csv` and `results_stage4_compare_vs_fixed_a1.csv`; rule: highest-objective archived auto policy, with the matched comparison taken against the generated fixed-$\alpha = 1.0$ baseline table.
- **Week 9.** Run `20260315_232052_week9_beta_vae`; source files: `beta_recon_summary.csv` and `beta_best_auto.csv`; rule: reconstruction summary and best safe policy reported separately for each $\beta \in \{1, 2, 4\}$.
- **Week 10 headline tuned-global result.** Run `20260319_214326_week10_dsprites_ood_only` inside `output_vae_tuned.zip`; source files: `tables/selected_comparison.csv`, `tables/auto_alpha_summary.csv`, `tables/fixed_alpha_sweep.csv`, and `tables/probe_eval.csv`; rule: among tuned-global reruns with the same protocol, select the run with the highest OOD auto-steering objective.
- **Week 10 supplementary local result.** Run `20260319_232434_week10_dsprites_ood_local_knn` inside `output_local_steering.zip`; source file: `tables/selected_comparison.csv`; rule: included for auditability as a preliminary local comparison, not as the headline week-10 stage in Table 1.

References

1. D. P. Kingma and M. Welling, "Auto-Encoding Variational Bayes," *International Conference on Learning Representations (ICLR)*, 2014.
2. I. Higgins, L. Matthey, A. Pal et al., "beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework," *International Conference on Learning Representations (ICLR)*, 2017.
3. Y. Bengio, A. Courville, and P. Vincent, "Representation Learning: A Review and New Perspectives," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(8):1798–1828, 2013.
4. F. Locatello, S. Bauer, M. Lucic et al., "Challenging Common Assumptions in the Unsupervised Learning of Disentangled Representations," *International Conference on Machine Learning (ICML)*, 2019.
5. L. Matthey, I. Higgins, D. Hassabis, and A. Lerchner, "dSprites: Disentanglement testing sprites dataset," 2017.

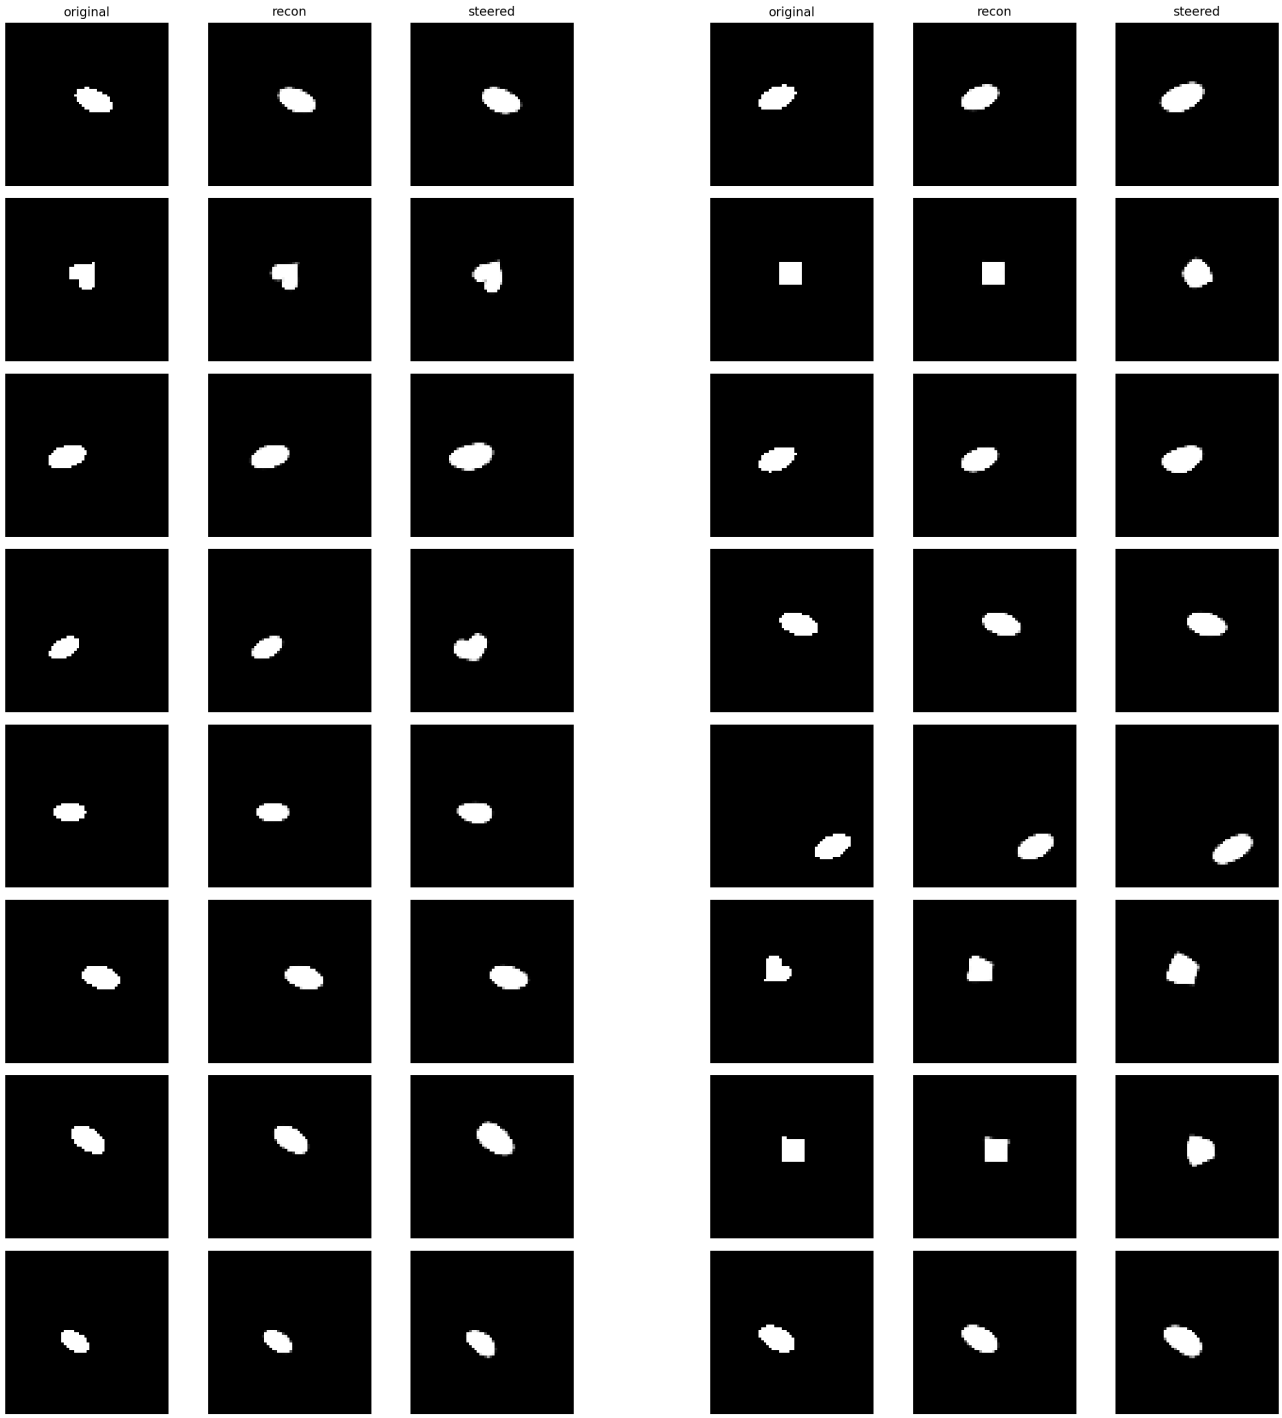

(a) ID preview (b) OOD preview

Figure 6: **Qualitative week-10 previews for tuned global dSprites steering.** The edit is visible but intentionally conservative in the safe regime; stronger fixed steps produce larger changes but noticeably worse preservation.